

Distributed URL Shortener

This project is a high-performance URL shortening service designed to handle heavy read/write traffic using a microservices-oriented architecture. By leveraging Redis for caching and Kafka for asynchronous event processing, the system achieves low redirect latency.

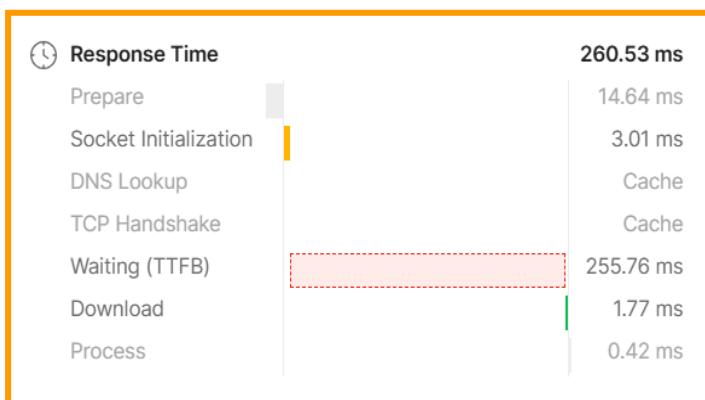
Key Technologies: * Backend: Node.js, Express.

- Database: MongoDB.
- Performance: Redis (Cache-Aside Pattern).
- Scalability: Apache Kafka & Zookeeper.
- Infrastructure: Nginx & Docker Compose

2. Problem Statement: The Latency Gap

In a standard URL shortener, every request requires a database lookup. During initial stress testing, my system exhibited high latency because every redirect was "blocking" while waiting for MongoDB to find the record and update the click count.

- **Observed Latency:** ~260ms.
- **The Goal:** Reduce latency to under 50ms and decouple analytics from the main request-response cycle.

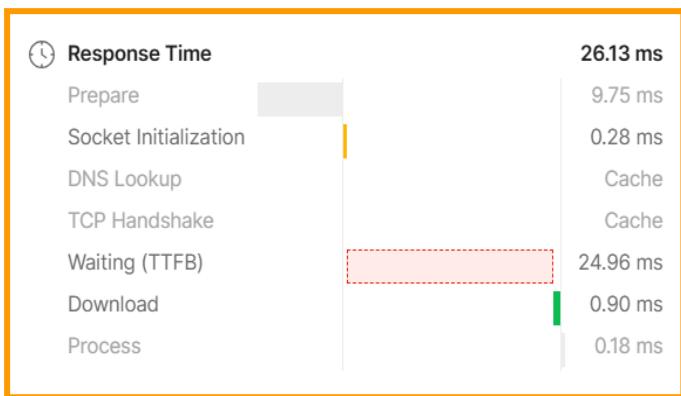


3. Architecture & Performance Engineering

a) Redis Caching (Cache-Aside Strategy)

I implemented Redis to store the mapping of `shortId` to `redirectURL`. This prevents redundant database queries for frequently accessed links.

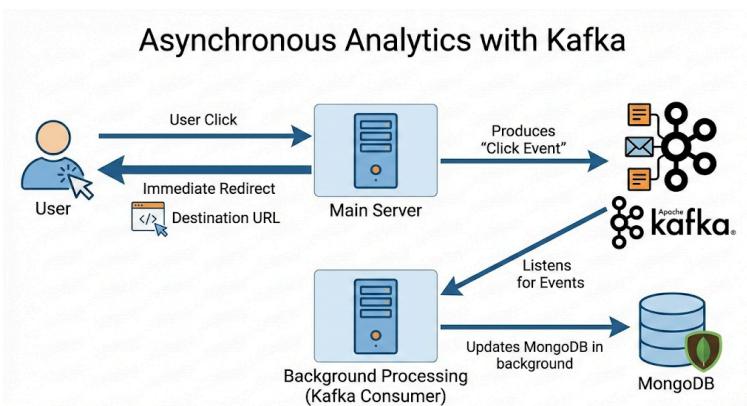
- Implementation: The backend checks Redis first; if data is missing, it fetches from MongoDB and "warms" the cache for 24 hours.
- Result: Redirect latency dropped to ~26ms, a 90% improvement in speed.



b). Asynchronous Analytics with Kafka

To ensure the redirect is never delayed by analytics logging, I integrated Apache Kafka.

- Process: The main server produces a "Click Event" to a Kafka topic and redirects the user immediately.
- Background Processing: A separate Kafka Consumer listens for these events and updates MongoDB in the background without affecting the user's experience.



4. Security: Distributed Rate Limiting

To protect the system from DDoS attacks and automated scraping, I developed a **Distributed Rate Limiter** using Redis.

- **Algorithm:** I utilized the **Token Bucket Algorithm** to allow for small traffic bursts while enforcing a strict upper limit (configured to 10 requests/minute per IP).
- **Distributed State:** Rate limits are keyed by User ID/IP in a centralized Redis cluster. This ensures that the security policy is enforced consistently across all **Docker containers**, preventing users from bypassing limits by hitting different backend instances.

The image contains two screenshots of the Postman application interface, both showing a POST request to `http://localhost/api/url`.

Screenshot 1 (Successful Request):

- The URL bar shows `http://localhost/api/url`.
- The method dropdown shows `POST`.
- The body content is `{ "url": "https://www.github.com" }`.
- The response status is `201 Created`.
- The response body is `{ "id": "gVq" }`.

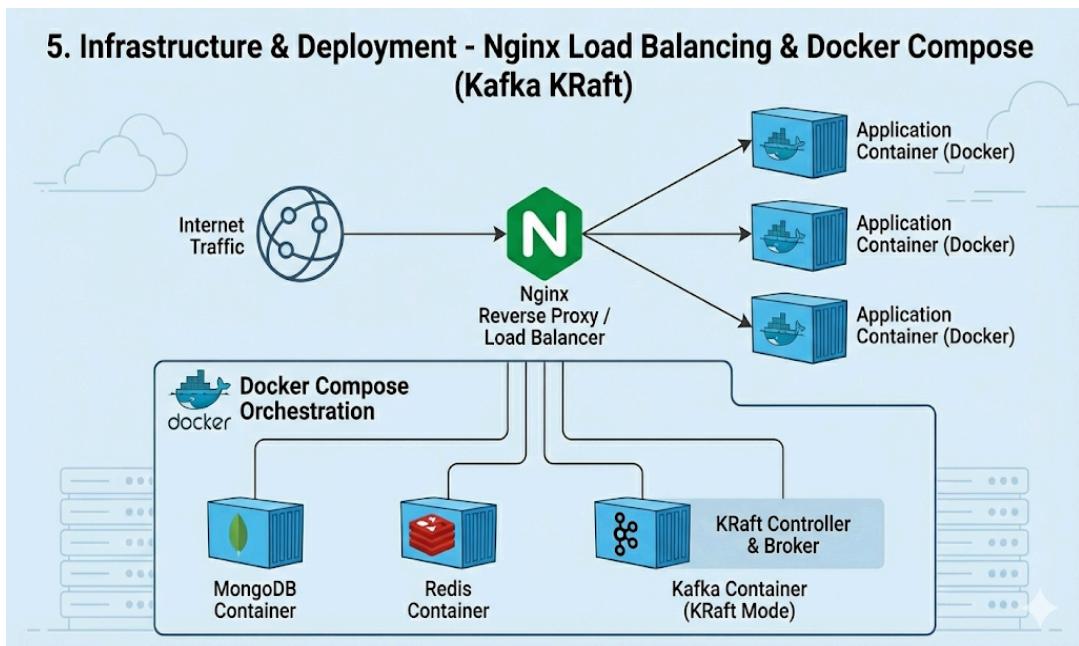
Screenshot 2 (Failed Request):

- The URL bar shows `http://localhost/api/url`.
- The method dropdown shows `POST`.
- The body content is `{ "url": "https://google.com" }`.
- The response status is `429 Too Many Requests`.
- The response body is `{ "error": "Too Many Requests. Please try again later." }`.

5. Infrastructure & Deployment

The entire system is containerized for consistent deployment across any environment.

- Load Balancing: Nginx acts as a Reverse Proxy, distributing incoming traffic.
- Orchestration: Docker Compose manages the lifecycle of the Application, MongoDB, Redis, Kafka .



6. Technical Challenges & Solutions

Challenge 1: Data Consistency & Atomic Updates

- **The Issue:** During initial integration, I encountered a data inconsistency bug where the database operation result (e.g., `{ modifiedCount: 1 }`) was inadvertently cached in Redis instead of the actual URL document. This caused subsequent redirect requests to fail.
- **The Solution:** I refactored the controller logic to utilize `findOneAndUpdate`. This atomic operation ensures the updated

document is returned immediately and cached correctly, maintaining data integrity between MongoDB and Redis.

Challenge 2: Distributed Message Durability

- **The Issue:** A critical requirement was ensuring that click analytics were never lost, even if a broker went offline during a traffic spike.
- **The Solution:** I configured the Kafka cluster using **KRaft** (Kafka Raft Metadata mode) instead of Zookeeper. This modern architecture simplifies the deployment while ensuring reliable leader election and message persistence across the cluster.

Challenge 3: Concurrency & Race Conditions

- **The Issue:** In a distributed system, simultaneous requests can lead to "Race Conditions" where the rate limiter counts incorrectly. If two requests read the counter at **9** at the exact same millisecond, both might increment to **10** and pass, breaking the strict limit.
- **The Solution:** I utilized **Redis Atomic Operations** (**INCR** and **EXPIRE**). Since Redis operates on a single-threaded event loop, it processes these operations sequentially. This guarantees 100% accurate counting and eliminates race conditions, even under high concurrency.